



The C++ Standard Template Library: Algorithms

Concepts, usage, and best practices

Author: Pavlos Sakoglou, M.Sc.

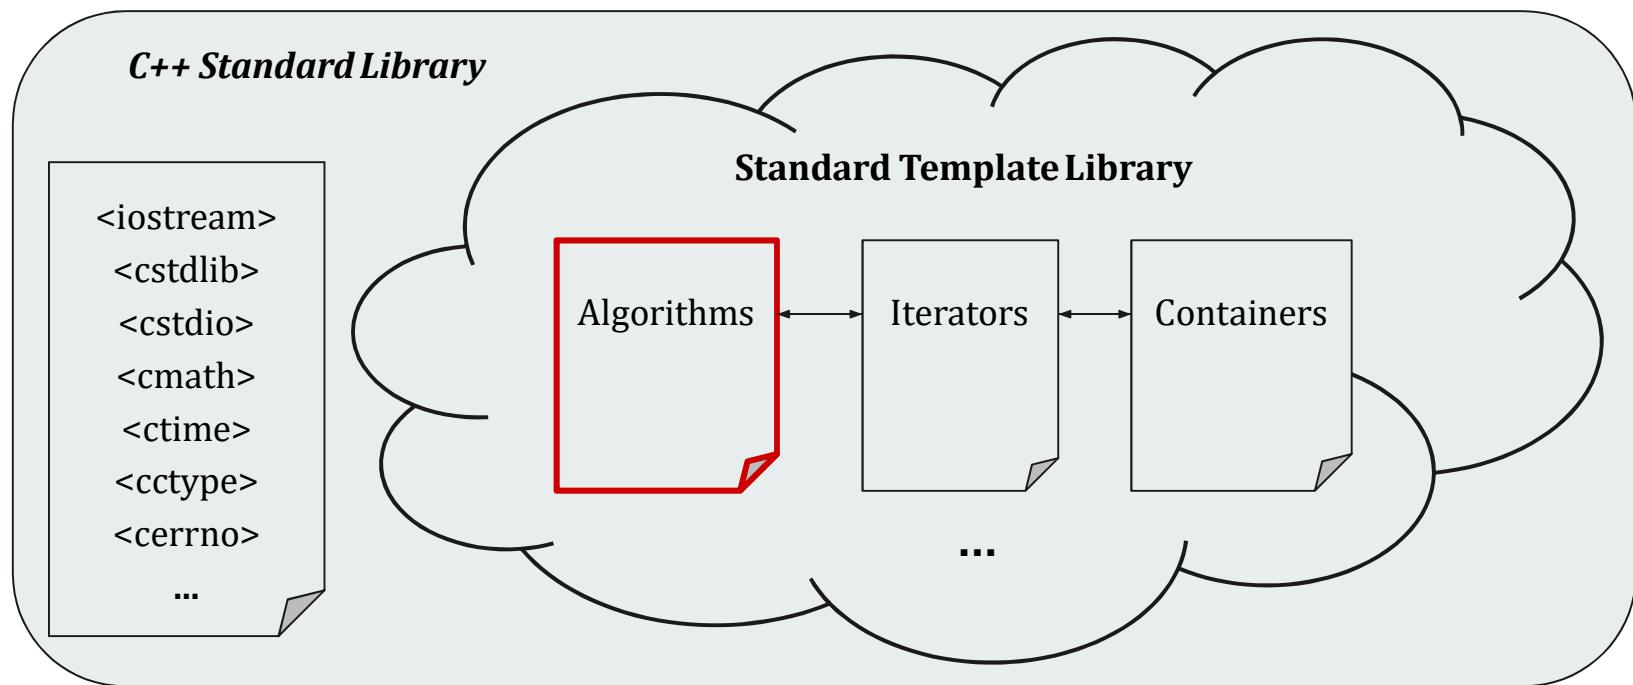
pavlos@superharmonic.tech / github.com/PavlosSakoglou

Outline



- **Brief review of STL**
- STL Algorithms overview
- Elementary STL algorithms
- Bugs and Pitfalls
- Performance in development
- References
- Q&A

Brief review of STL



Brief review of STL

- *STL Containers*: collection of most useful data structures
 - Sequence Containers
 - `std::vector`, `std::list`, `std::forward_list`, `std::deque`, `std::array`
 - Container Adaptors
 - `std::queue`, `std::stack`, `std::priority_queue`
 - Ordered/Unordered Associative Containers
 - `std::map`, `std::multimap`, `std::set`, `std::multiset`, `std::unordered_map`, `std::unordered_multimap`, `std::unordered_set`, `std::unordered_multiset`

More STL containers here: <http://en.cppreference.com/w/cpp/container>

Brief review of STL




- ***STL Iterators:*** overloaded pointer interface for STL containers
 - ***Random access:*** access elements at any arbitrary offset position
 - ***Bidirectional:*** access elements by moving in both directions
 - ***Forward:*** access elements by only moving forward
 - ***Contiguous*** (C++17): like random access but logically adjacent elements also adjacent in memory

More iterators: <http://en.cppreference.com/w/cpp/iterator>


Brief review of STL

➤ *STL Iterators*

```
std::vector<int> myVector = {1, 2, 3, 4, 5, 6};
```



```
std::list<double> myList = {1.1, 2.2, 3.5, 4.8, 5.7, 6.01};
```



Outline



- Brief review of STL
- **STL Algorithms overview**
- Elementary STL algorithms
- Bugs and Pitfalls
- Performance in development
- References
- Q&A

STL Algorithms overview



- What are STL algorithms?
 - Pre-built library of general purpose algorithms designed to solve specific problems i.e. sorting, searching, finding subsets, etc.
 - Free and standardized, used by thousands of developers daily
 - Declarative on syntax i.e. eliminate the need to write explicit loops
 - Iterate over some or all members of a *sequenced* STL container performing an operation on each element in turn
 - Designed by experts in numerical analysis and mathematics

STL Algorithms overview



- Why use STL algorithms?
 - Often more efficient than our own algorithms
 - Can leverage on the mechanics of STL containers
 - Cleaner code and clearly abstracted for several containers
 - Possible side effects contained inside the algorithm interface
 - Less likely to fail under non-obvious conditions
 - Intuitive identifiers, names, and interface
 - No need to reinvent the wheel

STL Algorithms overview



➤ *STL Algorithms (Part 1)*

- Non-modifying sequence operations
 - `std::for_each`, `std::count_if`, `std::search`, `std::find_if`, `std::equal`, ...
- Modifying sequence operations
 - `std::copy_if`, `std::reverse`, `std::replace_if`, `std::transform`, ...
- Sorting operations
 - `std::sort`, `std::is_sorted`, `std::partial_sort`, ...
- Numeric operations
 - `std::accumulate`, `std::inner_product`, `std::partial_sum`, `std::iota`, ...

STL Algorithms overview



➤ *STL Algorithms (Part 2)*

- Set operations
 - `std::merge`, `std::includes`, `std::set_union`, `std::inplace_merge`, ...
- Max/Min operations
 - `std::max_element`, `std::min_element`, `std::minimax_element`, ...
- Heap operations
 - `std::is_heap`, `std::make_heap`, `std::push_heap`, `std::sort_heap`, ...

More STL algorithms: <http://en.cppreference.com/w/cpp/algorithm>

Outline



- Brief review of STL
- STL Algorithms overview
- **Elementary STL algorithms**
- Bugs and Pitfalls
- Performance in development
- References
- Q&A

Elementary STL algorithms

- `std::for_each`
- Defined in `<algorithm>` http://en.cppreference.com/w/cpp/algorithm/for_each

```
template <class InputIter, class UnaryFunction>
UnaryFunction for_each(InputIter first, InputIter last, UnaryFunction f) {

    for (; first != last; ++first)
        f(*first);
    return f;
}
```

// Input: a range given by two iterators, and a policy functor to apply on that range

// Output: policy functor or void

// Complexity: if policy $f \sim O(f(n))$, then $(last - first) * O(f(n))$ i.e. *exactly (last - first) linear applications of f*

Elementary STL algorithms

- `std::for_each` example: print all even numbers in a container

// "print if even" policy

```
struct print_if_even {  
    void operator()(const int & number) {  
        if (number % 2 == 0) std::cout << number << " ";  
    }  
};
```

// ...


```
std::vector<int> myVector = { 1, 2, 3, 4, 5, 6, 7, 8, 9, 10 };
```

// Print all even numbers of vector

```
std::for_each(myVector.begin(), myVector.end(), print_if_even());
```

// Will apply the policy to each element and result to:

```
// { print_if_even()(1), print_if_even()(2), ..., print_if_even()(10) } = { 1, print(2), 3, print(4), ..., 9, print(10) }
```



Elementary STL algorithms

➤ `std::for_each` example: “print if even” with offset

```
// print all even values of starting at the middle of the vector
```

```
// ...
```

```
std::vector<int> myVector = { 1, 2, 3, 4, 5, 6, 7, 8, 9, 10 };
```

```
// Get the middle and end of container using const iterators
```

```
std::vector<int>::const_iterator middle = myVector.begin() + myVector.size() / 2;
```

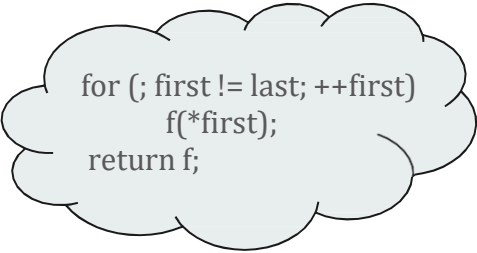
```
std::vector<int>::const_iterator end = myVector.end();
```

```
auto printing_evens = std::for_each(middle, end, print_if_even());
```

```
// What is 'auto'? What is the type of 'printing_evens'?
```

```
// What will the following line do?
```

```
// int k = 100; printing_evens(k);
```



```
for (; first != last; ++first)
    f(*first);
return f;
```

Elementary STL algorithms

- `std::transform`
- Defined in `<algorithm>` <http://en.cppreference.com/w/cpp/algorithm/transform>

```
template<class InputIter, class OutputIter, class UnaryFunction>
OutputIter transform(InputIter first, InputIter last, OutputIter d_first, UnaryFunction unary_f) {
    while (first != last) {
        *d_first++ = unary_f(*first++);
    }
    return d_first;
}
```

```
// Input: a range of a container of values, an output iterator to assign the transformed values, and an operation
// Output: output iterator + (last - first) i.e. output container's end()
// Complexity: if unary_f ~ O(f(n)) then (last - first) * O(f(n)) i.e. exactly last - first applications of unary_f
```

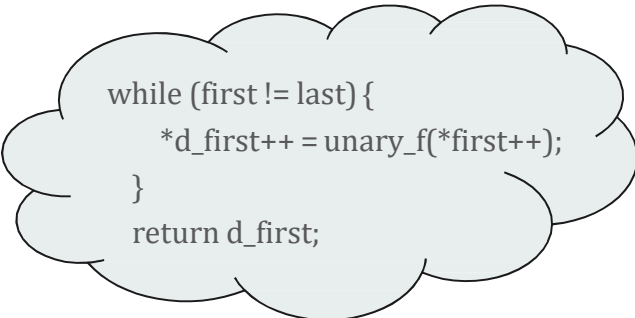

Elementary STL algorithms

- `std::transform` example: exponentiate an array

```
template <class NumericType>
struct exponentiate {
    NumericType operator()(NumericType & element)
    { return std::exp(element); }
};
```

```
// ...
std::array<double, 5> myArray = { 3.4, 2.5, 9.8, 12.01 };
```

```
// Transform the current array
std::transform(myArray.begin(), myArray.end(), myArray.begin(), exponentiate<double>());
```



```
while (first != last) {
    *d_first++ = unary_f(*first++);
}
return d_first;
```

Elementary STL algorithms

- `std::copy_if`
- Defined in `<algorithm>` <http://en.cppreference.com/w/cpp/algorithm/copy>

```
template <class InputIter, class OutputIter, class UnaryPredicate>
OutputIter copy_if(InputIter first, InputIter last, OutputIter d_first, UnaryPredicate pred) {
    while (first != last) {
        if (pred(*first)) *d_first++ = *first;
        first++;
    }
    return d_first;
}
```

```
// Input: a range of values to be copied, a start iterator to use for copying, and the predicate
// Output: the begin iterator of the new range
// Complexity: Linear  $O(\text{last} - \text{first})$  i.e. exactly last - first assignment operator calls
```

Elementary STL algorithms

- `std::copy_if` example: copy all negative numbers to a vector

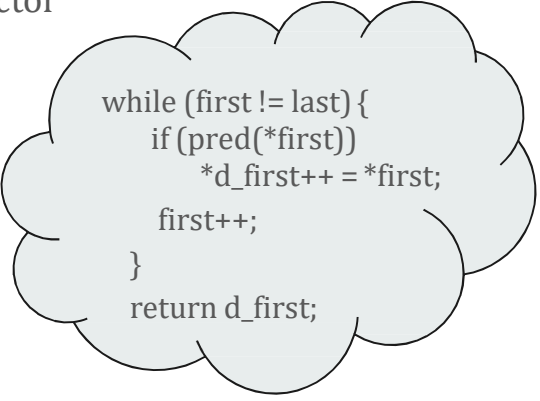
```
template <class NumericType>
struct is_negative {
public:
    bool operator()(const NumericType & value) { return value < 0; }
};

// ...
std::list<double> myList = { -1.1, 2.231, -3.23, 4.01, -9.1, -89.1, 12, 8.28 };
```

```
std::vector<double> negatives;
```

```
std::copy_if(myList.begin(), myList.end(), std::back_inserter(negatives), is_negative<double>());
```

```
// What is std::back_inserter? http://en.cppreference.com/w/cpp/iterator/back\_inserter
// copy_if does not do a push_back, only copies the value on the position where the iterator is
```



```
while (first != last) {
    if (pred(*first))
        *d_first++ = *first;
    first++;
}
return d_first;
```

Elementary STL algorithms

- `std::accumulate`
- Defined in `<numeric>` <http://en.cppreference.com/w/cpp/algorithm/accumulate>

```
template<class InputIter, class T>
T accumulate(InputIter first, InputIter last, T init) {

    for (; first != last; ++first)
        init = init + *first; // init is an offset
    return init;
}

// Input: a range of a container which values you want to accumulate (sum) and an initial value
// Output: the sum of values of the input range + the initial offset
// Complexity: linear O(last - first) i.e. exactly last - first addition operations
```

Elementary STL algorithms

➤ `std::accumulate` example: compute weekly investment earnings

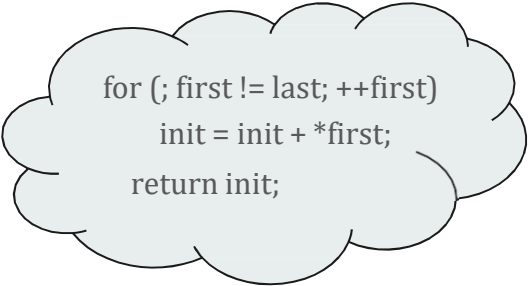
```
double initial_earnings = 5123.98;

// Weekly earnings
std::array<double, 7> earnings = {10.3, -12.01, -8.9, 5.1, -7.8, 12.0, 1.1};

// Compute the new earnings by accumulating the daily values
double new_earnings =
    std::accumulate(earnings.begin(), earnings.end(), initial_earnings);

if (new_earnings > initial_earnings)
    std::cout << "Good job!\n";
else
    std::cout << "We had losses!\n";

initial_earnings = new_earnings;
```



```
for (; first != last; ++first)
    init = init + *first;
return init;
```

Outline



- Brief review of STL
- STL Algorithms overview
- Elementary STL algorithms
- **Bugs and Pitfalls**
- Performance in development
- References
- Q&A

Bugs and Pitfalls



- Need to know involved STL features to use algorithms e.g. iterators
- You are not avoiding using pointers and their “buggy” behavior
- Newest algorithms won’t work with out-of-date compilers
- Doesn’t support native C++ code, only STL
- Easy to trigger run-time errors
- Not bug- or thread- safe

Bugs and Pitfalls



- Most common bugs when using STL algorithms:
 - Out-of-range input/output container iterators
 - Dereferencing *illegal* iterators implicitly
 - Poor container choices per algorithm
 - Container iterator const-correctness
 - Ignoring the algorithm's return type

Bugs and Pitfalls

➤ Out-of-range input/output iterators

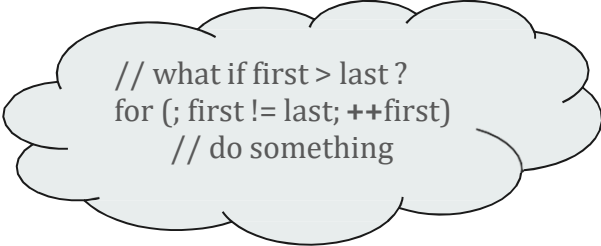
```
std::vector<int> myVector = {1, 2, 3, 4, 5};  
std::size_t m = myVector.size() + k; // (int) k > 0
```

```
std::for_each(myVector.begin(), myVector.end() - m, policy<int>()); // yikes!  
std::accumulate(myVector.begin() + m, myVector.end(), 0); // ouch!
```

// Error: Expression: vector iterator + offset out of range

```
std::accumulate(myVector.end(), myVector.begin(), 0); // crash!
```

// Error: Expression: invalid iterator range




```
// what if first > last?  
for (; first != last; ++first)  
    // do something
```

Bugs and Pitfalls

➤ Out-of-range input/output iterators

```
std::vector<int>::iterator start = myVector.begin();  
std::vector<int>::iterator end = myVector.end();  
int m = size + k; // (int) k > 0
```



// Error: Expression: vector iterator + offset out of range

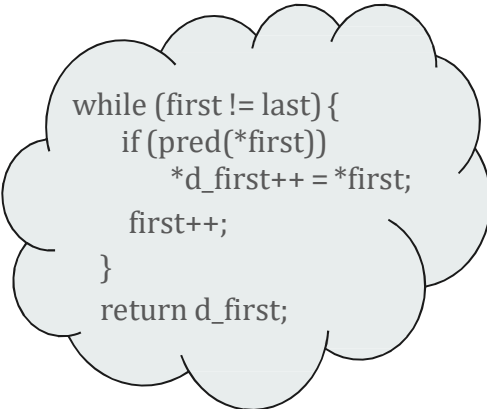
Bugs and Pitfalls

➤ Dereferencing *illegal* iterators

```
// Implicit illegal dereference
// ...
std::list<double> myList2 = { -1.1, 2.231, -3.23, 4.01, -9.1, -89.1, 12, 8.28 };
std::vector<double> negatives;
```

```
// The following line will throw an exception:
//     std::copy_if(myList2.begin(), myList2.end(), negatives.begin(), is_negative<double>());
```

```
// first = myList2.begin();
// last = myList2.end();
// d_first = negatives.begin(); // negatives is empty !
// ...
// *d_first++ = *first; // BUG: dereferencing uninitialized iterator !
// ...
// Error message: vector iterator not incrementable (why?)
```



```
while (first != last) {
    if (pred(*first))
        *d_first++ = *first;
    first++;
}
return d_first;
```

Bugs and Pitfalls

➤ Poor container choices and const-correctness

// **std::set**

```
std::set<int> mySet = { 2, 3, 5, 7, 9 };
```

```
// std::reverse(mySet.begin(), mySet.end()); // Compiler error !
```

```
// C3892: https://msdn.microsoft.com/en-us/library/cab7058b.aspx
```

```
// std::set is inherently constant. You can add and remove elements, but not modify/permute them
```

// **std::unordered_map**

```
std::unordered_map<int, int> myHash = { {1,1}, {2,2}, {3,3} };
```

```
// std::sort(myHash.begin(), myHash.end()); // Compiler errors !
```

```
// No comparison operator for std::unordered_map -- could not deduce type etc.
```

Bugs and Pitfalls

➤ Ignoring the return type

// Explicit illegal dereference

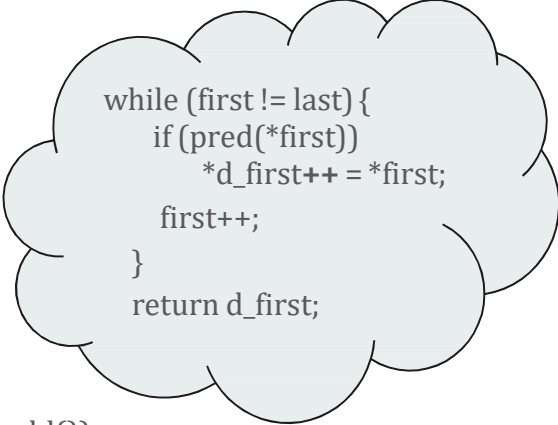
```
std::array<int, 10> intArray = {1, 2, 3, 4, 5, 6, 7, 8, 9, 10};  
std::array<int, 5> oddArray;
```

```
std::array<int, 5U>::iterator d_first =  
    std::copy_if(intArray.begin(), intArray.end(), oddArray.begin(), is_odd());
```

// After the last `*d_first++ = *first` assignment, the `++` operator returns and
// increments `d_first` to one position past the last element

// `std::cout << *d_first << std::endl;` // ouch!

// Error: Expression: array iterator not dereferencable



```
while (first != last) {  
    if (pred(*first))  
        *d_first++ = *first;  
    first++;  
}  
return d_first;
```

Outline



- Brief review of STL
- STL Algorithms overview
- Elementary STL algorithms
- Bugs and Pitfalls
- **Performance in development**
- References
- Q&A

Performance in development

- Problem 1: sort a vector of chars in increasing and decreasing order (part 1)

// Native C++ solution on github: https://github.com/PavlosSakoglou/STL-Algorithms-Tutorial-1/blob/master/Problem1_Native.cpp

```
template <class Type>
int partition(std::vector<Type> & myVector, int begin, int end) { /*...*/ }

template <class Type>
void myQuicksort(std::vector<Type> & myVector, int begin, int end) { /*...*/ }

// Sort
myQuicksort<char>(myVector, 0, size - 1); // O(n lg n)

// Copy values to increasing-order vector
for (const double & elem : myVector)
    increasing.push_back(elem);

// Reverse copy
auto begin = myVector.begin(); auto end = myVector.end(); int i = 1;
for (; begin != end; ++begin) decreasing.push_back(*(end - i++));
```

Performance in development

- Problem 1: sort a vector of chars in increasing and decreasing order

// STL solution on github: https://github.com/PavlosSakoglou/STL-Algorithms-Tutorial-1/blob/master/Problem1_STL.cpp

```
// STL Sort
std::sort(myVector.begin(), myVector.end()); // O(n lg n)
```

```
// Copy increasing
std::copy(myVector.begin(), myVector.end(), std::back_inserter(increasing)); // O(n)
```

```
// Copy decreasing
std::reverse_copy(myVector.begin(), myVector.end(), std::back_inserter(decreasing)); // O(n)
```


Performance in development

- Problem 2: computing basic and nth order statistics (part 1)

// Native C++ solution on github: https://github.com/PavlosSakoglou/STL-Algorithms-Tutorial-1/blob/master/Problem2_Native.cpp

```
// ...
template <class Type>
void myQuicksort(std::vector<Type> & myVector, int begin, int end) {/*...*/}

// Average
for (const double & elem : prices)
    daily_average += elem;
daily_average /= (double)size;

// Median
myQuicksort(prices, 0, size - 1);
daily_median = prices[size / 2 - 1];

// Copy max 5 prices
for (unsigned i = 0; i < 5; ++i)
    max_5_prices[i] = prices[size - 1 - i];
```

Performance in development

- Problem 2: computing basic and nth order statistics (part 2)

// Native C++ solution on github: https://github.com/PavlosSakoglou/STL-Algorithms-Tutorial-1/blob/master/Problem2_Native.cpp

```
// Max - Min prices
auto first = std::begin(prices);
auto last = std::end(prices);

if (first != last)
    min_price = max_price = *first++;

for (const double & elem : prices) {
    if (elem > max_price)
        max_price = elem;

    if (elem < min_price)
        min_price = elem;
}
```

Performance in development

➤ Problem 2: computing basic and nth order statistics

// STL solution on github: https://github.com/PavlosSakoglou/STL-Algorithms-Tutorial-1/blob/master/Problem2_STL.cpp

```
// Average
daily_average = std::accumulate(std::begin(prices), std::end(prices), 0.0) / (double)size;

// Median
std::partial_sort(std::begin(prices), std::begin(prices) + (size / 2) + 1, std::end(prices));
daily_median = prices[size / 2 - 1];

// Max - Min price
auto max_price_position = std::max_element(std::begin(prices), std::end(prices));
if (max_price_position != prices.end()) max_price = *max_price_position;

auto min_price_position = std::min_element(std::begin(prices), std::end(prices));
if (min_price_position != std::end(prices)) min_price = *min_price_position;

// Max 5 prices
std::nth_element(std::begin(prices), std::begin(prices) + 5, std::end(prices), decreasing_order<double>());
std::copy(std::begin(prices), std::begin(prices) + 5, std::back_inserter(max_5_prices));
```

Performance in development

➤ Problem 2: computing basic and nth order statistics

// BONUS: compute variance using `for_each` and a lambda expression vs C++11 range-based for loop

```
double variance = 0.0;

// STL
std::for_each(std::begin(prices), std::end(prices), [&](double elem) {
    variance += std::pow(elem - daily_average, 2);
});
```

```
// Native C++
for (const double & elem : prices)
    variance += std::pow(elem - daily_average, 2);
```

```
// Compute the unbiased variance
variance /= (double)(size - 1);
```

// `std::for_each` or range-based for-loop?

Outline



- Brief review of STL
- STL Algorithms overview
- Elementary STL algorithms
- Bugs and Pitfalls
- Performance in development
- **References**
- Q&A

References



- Thorough online documentation:
 - <http://en.cppreference.com/w/cpp/algorithm>
 - <http://www.cplusplus.com/reference/algorithm/>
 - <http://www.sgi.com/tech/stl/>
- Suggested open-source texts:
 - <http://stepanovpapers.com/STL/DOC.PDF>
 - <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2012/n3351.pdf>
- Presentation repository with all examples and codes:
 - <https://github.com/PavlosSakoglou/STL-Algorithms-Tutorial-1>

References



- What's next?
 - Advanced STL algorithms and iterators
 - C++ Lambda functions and STL function objects
 - Augmenting STL data structures and algorithms
 - STL algorithms and memory management

That is all!





The C++ Standard Template Library: Algorithms

Concepts, usage, and best practices

Thank you for your time!

Author: Pavlos Sakoglou, M.Sc.

pavlos@superharmonic.tech / github.com/PavlosSakoglou